

THE CMO GUIDE

Buying media when the execution layer runs itself

The execution is automated. The judgment isn't.

Published September 2026 · Updated October 2026

CONTENTS

Table of contents

01	Introduction	3
02	The new role of the CMO in 2026	4
03	The structural constraint	5
04	Why the execution layer is disappearing	6
05	The habit tax	7
06	What to show the board	8

07	Where CTV and retail media stand today	9
08	The agentic commerce shift	10
09	Zero-to-one partnership vs. platform selling	11
10	Mistakes to avoid	12
11	Questions before your next budget cycle	13

Most guides about AI and marketing are written for the platforms **selling** the AI. This one isn't.

Over the past eighteen months, the execution layer of media buying — bidding, pacing, creative rotation, audience refresh — has quietly become something a machine does better than a person watching a dashboard. Google changed how Target CPA works. Amazon and OpenAI are extending DSP campaigns into conversational surfaces. Every major platform has shipped its own version of an AI ad agent.

None of that is the hard part anymore.

The hard part is deciding what the machine should be doing in the first place — which audiences are worth chasing, which channels have actually moved with your customers, which "reasonable exceptions" from six months ago have quietly become your media plan. That's judgment. It doesn't automate, and it's not for sale inside a platform's ad manager.

This guide is for CMOs and marketing leaders who want a clear-eyed view of what's changed, what hasn't, and what to do about it before the next budget cycle — not a product pitch dressed up as thought leadership.

The old CMO job was built around managing inputs. What's left is the part that was **always** supposed to be the job.

The old scorecard was activity — how much got shipped, how many campaigns ran, how full the funnel looked. That job is disappearing, not because CMOs are being replaced, but because most of what used to require a person to manage now runs itself — freeing that time for the decisions only a person can actually make. Bid strategy adjusts in real time. Creative variants rotate on their own. Reporting dashboards refresh without anyone pulling a report.

- Deciding which audiences are actually worth reaching, not just which ones are easiest to target
- Knowing when a channel that worked eighteen months ago has quietly stopped working
- Owning the three or four decisions per quarter that actually move the P&L, instead of the twenty that just look like progress
- Translating media performance into something the board can act on — not a report, a decision

The CMOs who do well in this environment aren't the ones with the most AI tools running. They're the ones who've gotten sharper about the decisions a tool can't make for them, and quicker to move their own time toward the upstream commercial thinking that actually decides outcomes.

The winners of AI won't be the ones with the best model. **They'll be the ones already built to move at its speed.**

Most people who make this argument stop at mindset. The real constraint is structural.

Being ready for AI starts with an organization already aligned and sized to operate efficiently in this new environment — not one retrofitting a legacy structure onto new tools.

That's very difficult to do when a legacy org chart carries several times the headcount the workflows actually require. The org chart itself becomes the drag — independent of how willing any individual is to update their GTM, recognize old systems should improve, or let new data override old assumptions. You can have the right instincts and still be structurally unable to move.

None of this is an argument for fewer people. It's an argument for unlocking the people already there — less time lost managing dashboards, more time spent on the upstream commercial thinking that actually moves outcomes.

So the org that adapts fastest isn't the one with the newest stack. It's the one already lean enough, and trusted enough internally, to let systems update without needing to protect headcount or territory along the way.

Same models. Same access. The winners were never going to be decided by who had the API key — they'll be decided by who wasn't carrying the weight of a structure built for a world that already moved on.

This isn't a prediction. It's already happened in pieces most marketing teams haven't had the bandwidth to connect — connecting them isn't anyone's full-time job, which is exactly the problem.

On August 17, 2026, Google changed how budget-limited campaigns running Target CPA and Target ROAS behave. The target used to function like a soft ceiling — one a campaign could beat and quietly keep beating, sometimes for months. Now, for any campaign marked "limited by budget," the system works to deliver closer to the number you actually entered. Campaigns that were beating their targets don't stay there; costs quietly drift up to meet the number you set. Not every agency flagged it ahead of time; plenty explained it after the invoice.

That's a small example of a much bigger pattern. Amazon and OpenAI have already announced a formal partnership extending Amazon DSP campaigns into ChatGPT's conversational ad surface, currently piloting with select U.S. advertisers. Every major platform has shipped or is shipping its own AI ad agent: a system that takes a budget and a goal and runs the campaign with minimal human input.

An ad agent optimizing a campaign is not the same as a strategist deciding what campaign should exist.

The platforms building these agents have every incentive to blur that line, because the agent spends your budget inside their own auction. That's the part that didn't get automated. It's the part most media plans are quietly missing.

The most expensive line item in most marketing budgets isn't a bad decision. It's six months of reasonable ones.

A media plan gets built once, with real thought behind it, at the start of a fiscal year or a new leadership era. Then it drifts. A channel keeps its budget because pulling it means an uncomfortable conversation. A vendor relationship continues because switching feels riskier than staying. An audience segment keeps getting targeted because it used to convert, not because anyone checked whether it still does.

None of these are bad decisions in isolation. Compounded over six months, they add up to a media mix that made sense a while ago and hasn't been re-examined since.

Two channels illustrate this better than most right now: CTV and retail media. Both have moved faster than most media plans have kept pace with. The fix isn't more spend — it's a standing habit of re-testing the plan against where the audience has actually gone.

Most marketing decks brought to a board are built to demonstrate effort: campaign counts, channel breakdowns, a full backlog, reach and impression totals. A useful board update fits on three slides.

- **Return.** When does spend turn into profit? What's the ROAS trajectory, and is it moving the right direction?
- **Impact.** The total effect of marketing on the P&L — not just ad efficiency, but the levers that aren't ads at all: existing customer base, pricing and promotion, headcount.
- **Outlook.** Three priorities for the quarter, and what could go wrong — not a long list of internal goals and targets.

What doesn't belong in that room: team operations, upper-funnel metrics like reach, CPM, CTR and CPI, campaign counts, and twenty-page reports. This is also a useful filter for what your own media partner should be reporting to you.

These two channels get talked about constantly and understood accurately less often.

CTV inventory quality and targeting precision have both improved substantially in the past two years — the gap between "CTV" as a line item and CTV as addressable, accountable media has mostly closed. A lot of plans still buy it the way linear TV used to be bought: broad, reach-based, hard to measure.

Retail media has moved even faster, in a different direction — it's increasingly where purchase-intent signal lives, not just where ad space gets sold. Brands treating it as a bottom-funnel afterthought are missing that these platforms now carry some of the best first-party signal available anywhere.

Both channels have moved further, faster, than most brands' media plans have.

A newer development, worth understanding even if it's not urgent yet: AI platforms are building the infrastructure to let their own systems discover and transact on a brand's behalf. OpenAI's Agentic Commerce Protocol and Google's Universal Commerce Protocol (both live, though still rolling out to a limited set of partners) are early moves toward a world where a consumer's AI assistant browses, compares, and potentially checks out — without visiting a website or seeing a traditional ad at all.

This isn't a channel yet in the way search or social are channels. It's closer to an emerging requirement: making sure product and catalog data is structured in a way these systems can read, trust, and act on. Brands that get this right early will show up correctly when AI-driven discovery becomes a meaningful share of commerce.

Know whether your product data is ready — and who's tracking this closely enough to tell you when it stops being optional.

There's a distinction worth being explicit about: the rep at a major retail media or DSP platform selling you ad inventory has a fundamentally different job than a media partner working on your behalf.

A platform rep is good at what the role is built for — keeping existing spend on their platform performing well. What the role isn't built for is the zero-to-one work: a fully custom program, built from nothing, with no guarantee it pays off — for one advertiser out of a whole book of accounts. That's not a knock on the rep. It's a reasonable outcome of the economics of the job. Most agency fee models run on the same logic — built around managing and executing a budget efficiently, not around spending unbilled time asking whether the budget belongs there in the first place.

No platform quotas or volume commitments to protect. Every cost laid out up front.

- **Expecting platform-agnostic advice from a platform rep.** They're genuinely good at making their own inventory perform — that's the job, and it's a real skill. What it isn't is unbiased: they're not positioned to tell you when a different platform, or no platform at all, is the better call. Not a knock on the person; just not what the role is for.
- **Letting "reasonable exceptions" become the plan.** Individually defensible decisions, left unexamined for six months, become a media mix nobody actually chose.
- **Reporting activity instead of payback.** If the update is full of reach and campaign counts, it's answering the wrong question.
- **Assuming automation removed the need for judgment.** The execution layer running itself means the plan matters more, not less.
- **Overlooking how the fee model shapes advice.** An agency built around execution seat licensing and media commission has less structural reason to recommend spending less, or moving to a channel outside its own toolset.
- **Waiting for a formal audit instead of building a standing habit.** By the time a plan gets a full review, it's usually been drifting for a year.

- 01 When was the current media mix last rebuilt from scratch, versus incrementally adjusted?
- 02 Which channels were chosen because they're working today, versus because they worked eighteen months ago?
- 03 Does your CTV and retail media strategy reflect where those channels actually stand now?
- 04 Is your reporting built around payback and profit, or around activity and reach?
- 05 Is your product data ready for AI-driven discovery, and who's tracking that for you?
- 06 If your agency's fee model rewards keeping spend where it already sits, who's telling you when a different platform is the right call?

If more than one or two of these don't have a confident answer, that's worth a conversation before the budget locks.

ABOUT ORBIS TEAM

Orbis Team is an independent programmatic, CTV, and retail media practice, built by practitioners with 35+ years of combined experience buying paid media for national and regional brands. This guide reflects how we already work: no platform quotas or volume commitments, and every cost laid out before anything runs.

The execution is automated. The judgment isn't.

orbisteam.com